

Kadwasra Ravi Kumar

Data & AI Engineer

Email: kadwasra619@gmail.com

LinkedIn: [linkedin.com/in/kadwasra-ravi-kumar-7b3b8318b](https://www.linkedin.com/in/kadwasra-ravi-kumar-7b3b8318b)

GitHub: github.com/dxdmikey

Portfolio: kadwasra.vercel.app

Professional Summary

Data & AI Engineer with 3+ years of experience designing scalable data platforms, lakehouse architectures, and enterprise ETL solutions on Azure. Currently building a self-serve data platform with an integrated AI agent for a construction engineering enterprise. Experienced in Azure Databricks, Synapse Analytics, Microsoft Fabric, PySpark, and SQL, and in building RAG-based AI agents on Azure. Proven track record of migrating large-scale data platforms, optimizing processing performance by up to 35%, and delivering reliable analytics-ready datasets. Ships full-stack applications end to end with AI-assisted development.

Experience

Navayuga Engineering Company Ltd, Hyderabad, India

Full-stack Developer & Data Engineer | Sep 2026 – Present

- Co-building, in a two-person Digital Intelligence & Transformation team, a self-serve data platform that lets non-technical users configure data sources, apply transformations, and build reports.
- Integrating an enterprise ERP with 2,400+ tables, a fuel-management API, fleet telematics, IoT, and truck data into a unified Azure data platform.
- Developing the reporting layer on Apache ECharts, enabling business teams to create and change reports without engineering support.
- Building an AI agent that answers data questions using retrieval-augmented generation (RAG) and generates charts from natural-language prompts.

WinWire Technologies, Hyderabad, India

Data Engineer SDE | May 2025 – Aug 2026

- Designed proof-of-concept data platforms using Microsoft Fabric, Azure Databricks, and AWS, accelerating client onboarding and architecture validation.
- Developed metadata-driven ingestion frameworks using Azure Data Factory and Synapse, reducing pipeline development effort by 40%.
- Implemented Medallion lakehouse architecture (Bronze–Silver–Gold), improving analytical query performance by 30%.
- Enhanced production monitoring, alerting, and deployment processes, improving platform reliability and reducing operational overhead.

Project– Commercial Data Platform (AMS Support) | Oct 2025 – Aug 2026

- Managed and optimized 25+ production pipelines processing multi-source commercial data from EDI, SFTP, and retail systems.
- Reduced recurring pipeline failures by 30% through root-cause analysis and proactive data quality controls.
- Improved SQL and PySpark processing performance by 20–35% through query tuning and transformation optimization.
- Built validation frameworks across source, staging, curated, and Synapse layers, ensuring data accuracy and consistency.
- Reduced incident resolution time (MTTR) by 25% through enhanced logging, monitoring, and alerting strategies.

Data Engineer SDT | Nov 2023 – Apr 2025

- Developed distributed data pipelines processing 50M+ records across enterprise healthcare systems.
- Delivered scalable ETL workflows using PySpark, SQL, Azure Synapse, and Databricks for large-scale analytical workloads.
- Supported cloud migration and data integration initiatives focused on performance, reliability, and data quality.

Project– Healthcare Data Migration (Apr 2024 – Apr 2025)

- Migrated 400+ Oracle tables to Azure Synapse and Databricks with 99.7% data accuracy.
- Developed 45+ Synapse pipelines and 20+ PySpark transformation frameworks supporting enterprise-scale migration.
- Implemented Delta Lake-based Medallion architecture, improving processing performance by 35%.
- Designed automated reconciliation and validation frameworks to ensure schema consistency and data integrity.
- Automated orchestration and alerting processes, significantly reducing manual operational effort.

Project– Healthcare Data Integration (Jan 2024 – Apr 2024)

- Engineered automated ingestion pipelines processing 30K+ daily healthcare and supply-chain records.
- Implemented incremental loading frameworks, reducing data refresh SLAs from 24 hours to under 4 hours.
- Improved pipeline stability by 30% through metadata-driven transformation design.
- Delivered analytics-ready datasets consumed by BI and reporting teams.

Education

Vellore Institute of Technology (VIT), Bhopal, India

B.Tech in Computer Science and Engineering (Specialization in AI & ML) (Jun 2019 – Jun 2023)

Final CGPA: 8.2 / 10

Certifications

AI-102: Azure AI Engineer Associate – Feb 2026

DP-100: Azure Data Scientist Associate – Dec 2025

DP-700: Microsoft Fabric Data Engineer Associate – Aug 2025

Databricks Data Engineer Associate – Mar 2024

Skills

Programming & Processing: Python, PySpark, SQL (Advanced), Spark SQL

Data Platforms: Azure Databricks, Azure Synapse, Microsoft Fabric, Delta Lake

Data Engineering: ETL/ELT Pipelines, Medallion Architecture, Incremental Processing, Data Modeling (Star, Snowflake, SCD Type 2)

Cloud & Storage: Azure Data Factory, ADLS Gen2, Azure SQL, AWS (S3, Glue, Redshift)

AI & Agents: Retrieval-Augmented Generation (RAG), AI Agents, Azure OpenAI, Prompt Engineering

Full-stack (AI-assisted): Next.js, React, REST APIs, Claude Code, Vercel

Orchestration & Optimization: Airflow, Pipeline Automation, Partitioning, Performance Tuning, Monitoring & Alerting

Analytics & Visualization: Apache ECharts, Power BI, Self-serve Reporting

Data Engineering Practice Projects

1. AWS Data Warehouse Platform (S3 + Glue + Redshift)

- Developed ETL pipelines using AWS Glue and PySpark to process data from Amazon S3 into Redshift.
- Designed dimensional warehouse tables and implemented incremental loading strategies for efficient data processing.
- Optimized Redshift queries using distribution and sort keys, improving analytical query performance.

2. Enterprise Data & Analytics Platform (ADF + ADLS Gen2 + Claude + Python)

- Built a configurable ingestion framework for Azure SQL, REST API and FTP/SFTP sources, orchestrated with Azure Data Factory into ADLS Gen2 (Bronze) and Azure SQL targets, with run tracking and duplicate handling.
- Designed a metadata-driven backend with FastAPI, PostgreSQL and SQLAlchemy/Alembic, including RBAC, hierarchical business-unit access and configurable workflow engines.
- Developed a Query Editor and an ECharts-based Dashboard Builder for dataset discovery, querying and interactive dashboards, backed by 500+ automated backend tests.

Languages: English, Telugu, Hindi